

## Experiment Number

Name ::	Rishabh Anand	UID ::	19BCS4525
Branch ::	CSE - IoT	Sec/Grp ::	1/A
Semester ::	6 <sup>th</sup>	Date ::	14 <sup>th</sup> Mar, 2022
Subject ::	ML Lab	CODE ::	CSD-386

### 1. Aim :

To implement the logistic regression on the given data set and tune the model with the appropriate selection of solver.

- DATA-SET LINK: <https://www.kaggle.com/dinaouahbi/churndatacsv>

### 2. Task :

1. Customer churn with Logistic Regression: A telecommunications company is concerned about the number of customers leaving their land-line business for cable competitors. They need to understand who is leaving. Imagine that you're an analyst at this company and you have to find out who is leaving and why.

### 3. Algorithm :

- Import required Libraries
- Load the Telco Churn Data
- Data pre-processing and selection
- Training and Testing
- Confussion Matrix

## 4. Code :

File Edit View Insert Cell Kernel Widgets Help

Not Tracked Python 3 (ipykernel)

### Customer churn with Logistic Regression

A telecommunications company is concerned about the number of customers leaving their land-line business for cable competitors. They need to understand who is leaving. Imagine that you're an analyst at this company and you have to find out who is leaving and why.

Lets first import required libraries:

```
In [1]: import pandas as pd
import pylab as pl
import numpy as np
import scipy.optimize as opt
from sklearn import preprocessing

import matplotlib.pyplot as plt
```

**Load the Telco Churn data**

Telco Churn is a hypothetical data file that concerns a telecommunications company's efforts to reduce turnover in its customer base.

Each case corresponds to a separate customer and it records various demographic and service usage information.

```
In [2]: churn_df = pd.read_csv("../DataSets/ChurnData.csv")
churn_df.head()
```

	tenure	age	address	income	ed	employ	equip	calcard	wireless	longmon	...	pager	internet	callwait	confer	ebill
0	11.0	33.0	7.0	136.0	5.0	5.0	0.0	1.0	1.0	4.40	...	1.0	0.0	1.0	1.0	0.0
1	33.0	33.0	12.0	33.0	2.0	0.0	0.0	0.0	0.0	9.45	...	0.0	0.0	0.0	0.0	0.0
2	23.0	30.0	9.0	30.0	1.0	2.0	0.0	0.0	0.0	6.30	...	0.0	0.0	0.0	1.0	0.0
3	38.0	35.0	5.0	76.0	2.0	10.0	1.0	1.0	1.0	6.05	...	1.0	1.0	1.0	1.0	1.0

File Edit View Insert Cell Kernel Widgets Help

Not Tracked Python 3 (ipykernel)

```
In [3]: churn_df.isnull().sum()
```

```
tenure      0
age          0
address      0
income       0
ed           0
employ       0
equip        0
calcard      0
wireless     0
longmon      0
tollmon      0
equipmon     0
cardmon      0
wiremon      0
longten      0
tollten      0
cardten      0
voice        0
pager        0
internet     0
callwait     0
confer       0
ebill        0
loglong      0
logtoll      0
lninc        0
custcat      0
churn        0
dtype: int64
```

```
In [4]: churn_df.info
```

```
<bound method DataFrame.info of      tenure  age  address  income  ed  employ  equip  calcard  wireless \
0      11.0  33.0    7.0    136.0  5.0    5.0    0.0    1.0    1.0
1      33.0  33.0   12.0    33.0  2.0    0.0    0.0    0.0    0.0
2      23.0  30.0    9.0    30.0  1.0    2.0    0.0    0.0    0.0
3      38.0  35.0    5.0    76.0  2.0   10.0    1.0    1.0    1.0
4       7.0  35.0   14.0    80.0  2.0   15.0    0.0    1.0    0.0
5      68.0  52.0   17.0   120.0  1.0   24.0    0.0    1.0    0.0
6      42.0  40.0    7.0    37.0  2.0    8.0    1.0    1.0    1.0
7       9.0  21.0    1.0    17.0  2.0    2.0    0.0    0.0    0.0
8      25.0  50.0   25.0   110.0  3.0   20.0    0.0    1.0    0.0
```

File Edit View Insert Cell Kernel Widgets Help
Not Trusted Python 3 (ipykernel)

In [4]: churn\_df.info

```
<bound method DataFrame.info of      tenure  age  address  income  ed  employ  equip  callcard  wireless  \
0      11.0  33.0    7.0   136.0    5.0    5.0    0.0    1.0    1.0
1      33.0  33.0   12.0    33.0    2.0    0.0    0.0    0.0    0.0
2      23.0  30.0    9.0    30.0    1.0    2.0    0.0    0.0    0.0
3      38.0  35.0    5.0    76.0    2.0   10.0    1.0    1.0    1.0
4       7.0  35.0   14.0    80.0    2.0   15.0    0.0    1.0    0.0
5      68.0  52.0   17.0   120.0    1.0   24.0    0.0    1.0    0.0
6      42.0  40.0    7.0    37.0    2.0    8.0    1.0    1.0    1.0
7       9.0  21.0    1.0    17.0    2.0    2.0    0.0    0.0    0.0
8      35.0  50.0   26.0   140.0    2.0   21.0    0.0    1.0    0.0
9      49.0  51.0   27.0    63.0    4.0   19.0    0.0    1.0    0.0
10     56.0  52.0   28.0    49.0    2.0   12.0    0.0    1.0    0.0
11     47.0  40.0   16.0   127.0    4.0   12.0    1.0    1.0    0.0
12     56.0  50.0    1.0    80.0    2.0   24.0    0.0    1.0    1.0
13     69.0  51.0   11.0   438.0    4.0   23.0    1.0    1.0    0.0
14     16.0  27.0    5.0    37.0    3.0    5.0    0.0    0.0    0.0
15      4.0  35.0   16.0   161.0    5.0    6.0    1.0    0.0    1.0
16     27.0  51.0    3.0    80.0    5.0   11.0    1.0    0.0    0.0
17     52.0  61.0    3.0    53.0    5.0    1.0    1.0    1.0    1.0
18     64.0  25.0    4.0    76.0    3.0    2.0    1.0    1.0    0.0
19     12.0  24.0    2.0    19.0    1.0    0.0    0.0    1.0    0.0
20     35.0  61.0   23.0    41.0    2.0   11.0    0.0    1.0    0.0
21     13.0  54.0    2.0    31.0    4.0    2.0    0.0    0.0    0.0
22     16.0  17.0    7.0    36.0    4.0    1.0    0.0    0.0    0.0
```

In [5]: churn\_df.shape

```
(200, 28)
```

### Data pre-processing and selection

Lets select some features for the modeling. Also we change the target data type to be integer.

In [6]:

```
churn_df = churn_df[['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip', 'callcard',
churn_df['churn'] = churn_df['churn'].astype('int')
churn_df
```

File Edit View Insert Cell Kernel Widgets Help
Not Trusted Python 3 (ipykernel)

In [6]:

```
churn_df = churn_df[['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip', 'callcard',
churn_df['churn'] = churn_df['churn'].astype('int')
churn_df
```

```
/home/fenis/.condahome/envs/Uni/lib/python3.7/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy
```

	tenure	age	address	income	ed	employ	equip	callcard	wireless	churn
0	11.0	33.0	7.0	136.0	5.0	5.0	0.0	1.0	1.0	1
1	33.0	33.0	12.0	33.0	2.0	0.0	0.0	0.0	0.0	1
2	23.0	30.0	9.0	30.0	1.0	2.0	0.0	0.0	0.0	0
3	38.0	35.0	5.0	76.0	2.0	10.0	1.0	1.0	1.0	0
4	7.0	35.0	14.0	80.0	2.0	15.0	0.0	1.0	0.0	0
5	68.0	52.0	17.0	120.0	1.0	24.0	0.0	1.0	0.0	0
6	42.0	40.0	7.0	37.0	2.0	8.0	1.0	1.0	1.0	0
7	9.0	21.0	1.0	17.0	2.0	2.0	0.0	0.0	0.0	0
8	35.0	50.0	26.0	140.0	2.0	21.0	0.0	1.0	0.0	0
9	49.0	51.0	27.0	63.0	4.0	19.0	0.0	1.0	0.0	0
10	56.0	52.0	28.0	49.0	2.0	12.0	0.0	1.0	0.0	0
11	47.0	40.0	16.0	127.0	4.0	12.0	1.0	1.0	0.0	0
12	56.0	50.0	1.0	80.0	2.0	24.0	0.0	1.0	1.0	0
13	69.0	51.0	11.0	438.0	4.0	23.0	1.0	1.0	0.0	0
14	16.0	27.0	5.0	37.0	3.0	5.0	0.0	0.0	0.0	0
15	4.0	35.0	16.0	161.0	5.0	6.0	1.0	0.0	1.0	1
16	27.0	51.0	3.0	80.0	5.0	11.0	1.0	0.0	0.0	0
17	52.0	61.0	3.0	53.0	5.0	1.0	1.0	1.0	1.0	0
18	64.0	25.0	4.0	76.0	3.0	2.0	1.0	1.0	0.0	0
19	12.0	24.0	2.0	19.0	1.0	0.0	0.0	1.0	0.0	1
20	35.0	61.0	23.0	41.0	2.0	11.0	0.0	1.0	0.0	0
21	13.0	54.0	2.0	31.0	4.0	2.0	0.0	0.0	0.0	0

File Edit View Insert Cell Kernel Widgets Help
Python 3 (ipykernel)

Selecting X and Y from the dataset

```

In [7]: X = np.asarray(churn_df[['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip']])
print(X[0:5])
y = np.asarray(churn_df['churn'])
print(y [0:5])

[[ 11.  33.   7. 136.   5.   5.   0.]
 [ 33.  33.  12.  33.   2.   0.   0.]
 [ 23.  30.   9.  30.   1.   2.   0.]
 [ 38.  35.   5.  76.   2.  10.   1.]
 [  7.  35.  14.  80.   2.  15.   0.]]
[1 1 0 0 0]

```

```

In [8]: print(y [0:100])

[1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0 0 0 0 0 0 1 0 0 1
 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0
 0 0 0 1 0 1 0 0 1 1 1 0 0 0 1 0 1 0 1 0 0 0 1 1 1 0 0]

```

Normalizing the dataset\*\*

```

In [9]: from sklearn import preprocessing

X = preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]

array([[ -1.13518441, -0.62595491, -0.4588971 ,  0.4751423 ,  1.6961288 ,
        -0.58477841, -0.85972695],
       [-0.11604313, -0.62595491,  0.03454064, -0.32886061, -0.6433592 ,
        -1.14437497, -0.85972695],
       [-0.57928917, -0.85594447, -0.261522 , -0.35227817, -1.42318853,
        -0.92053635, -0.85972695],
       [ 0.11557989, -0.47262854, -0.65627219,  0.00679109, -0.6433592 ,
        -0.02518185,  1.16316 ],
       [-1.32048283, -0.47262854,  0.23191574,  0.03801451, -0.6433592 ,
        0.53441472, -0.85972695]])

```

Training and Testing

```

In [10]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, random_state=4)
print ('Train set:', X_train.shape,  y_train.shape)
print ('Test set:', X_test.shape,  y_test.shape)

Train set: (160, 7) (160,)
Test set: (40, 7) (40,)

```

```

In [11]: y_test

array([0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0])

```

This function implements logistic regression and can use different numerical optimizers to find parameters, including 'newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga' solvers. There are pros and cons of these optimizers. The version of Logistic Regression in Scikit-learn, support regularization. Regularization is a technique used to solve the overfitting problem in machine learning models.

File
Edit
View
Insert
Cell
Kernel
Widgets
Help
Not Tracked
Python 3 (ipykernel)

Training and Testing

```

In [10]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)
print('Train set:', X_train.shape, y_train.shape)
print('Test set:', X_test.shape, y_test.shape)

Train set: (160, 7) (160,)
Test set: (40, 7) (40,)

```

```

In [11]: y_test

array([0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0])

```

This function implements logistic regression and can use different numerical optimizers to find parameters, including 'newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga' solvers. There are pros and cons of these optimizers. The version of Logistic Regression in Scikit-learn, support regularization. Regularization is a technique used to solve the overfitting problem in machine learning models. C parameter indicates inverse of regularization strength which must be a positive float. Smaller values specify stronger regularization.

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)

```

In [12]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
LR = LogisticRegression(C=0.001, solver='liblinear').fit(X_train,y_train)
LR

/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:35: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations

```

File
Edit
View
Insert
Cell
Kernel
Widgets
Help
Not Tracked
Python 3 (ipykernel)

```

In [12]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
LR = LogisticRegression(C=0.001, solver='liblinear').fit(X_train,y_train)
LR

/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:35: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
eps=np.finfo(np.float).eps,
/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:597: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
eps=np.finfo(np.float).eps, copy_X=True, fit_path=True,
/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:836: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
eps=np.finfo(np.float).eps, copy_X=True, fit_path=True,
/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:862: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
eps=np.finfo(np.float).eps, positive=False):
/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:1074: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
max_n_alphas=1000, n_jobs=1, eps=np.finfo(np.float).eps,
/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:1306: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
max_n_alphas=1000, n_jobs=1, eps=np.finfo(np.float).eps,
/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/least_angle.py:1442: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
eps=np.finfo(np.float).eps, copy_X=True, positive=False):
/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/randomized_l1.py:152: DeprecationWarning: 'np.float' is a deprecated alias for the builtin 'float'. To silence this warning, use 'float' by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use 'np.float64' here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations

```

File Edit View Insert Cell Kernel Widgets Help
Not Tracked Python 3 (ipykernel)

```

LogisticRegression(C=0.001, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)

```

Testing

In [13]:

```

yhat = LR.predict(X_test)
yhat

```

```

/home/fenris/.condahome/envs/Uni/lib/python3.7/site-packages/sklearn/linear_model/base.py:326: DeprecationWarning: 'np.int' is
a deprecated alias for the builtin 'int'. To silence this warning, use 'int' by itself. Doing this will not modify any behavio
r and is safe. When replacing 'np.int', you may wish to use e.g. 'np.int64' or 'np.int32' to specify the precision. If you wis
h to review your current use, check the release note link for additional information.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
indices = (scores > 0).astype(np.int)

array([0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0,
       0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1])

```

predict\_proba returns estimates for all classes, ordered by the label of classes. So, the first column is the probability of class 1, P(Y=1|X), and second column is probability of class 0, P(Y=0|X):

In [14]:

```

yhat_prob = LR.predict_proba(X_test)
yhat_prob

```

```

array([[0.5055063 , 0.4944937 ],
       [0.51347317, 0.48652683],
       [0.50602462, 0.49397538],
       [0.52000045, 0.47999955],
       [0.50661998, 0.49338002],
       [0.50673336, 0.49326664],
       [0.49839547, 0.50160453],
       [0.51405871, 0.48594129],
       [0.48122787, 0.51877213],
       [0.51944971, 0.48055029],
       [0.51298723, 0.48701277],
       [0.52293885, 0.47706115],
       [0.48984113, 0.51015887],

```

File Edit View Insert Cell Kernel Widgets Help

Not Tracked Python 3 (ipykernel)

In [14]:

```

yhat_prob = LR.predict_proba(X_test)
yhat_prob

```

```

array([[0.5055063 , 0.4944937 ],
       [0.51347317, 0.48652683],
       [0.50602462, 0.49397538],
       [0.52000045, 0.47999955],
       [0.50661998, 0.49338002],
       [0.50673336, 0.49326664],
       [0.49839547, 0.50160453],
       [0.51405871, 0.48594129],
       [0.48122787, 0.51877213],
       [0.51944971, 0.48055029],
       [0.51298723, 0.48701277],
       [0.52293885, 0.47706115],
       [0.48984113, 0.51015887],
       [0.48407622, 0.51592378],
       [0.52861954, 0.47138046],
       [0.5092273 , 0.4907727 ],
       [0.50191981, 0.49808099],
       [0.49233906, 0.50766094],
       [0.49487978, 0.50512022],
       [0.49864676, 0.50135324],
       [0.51699424, 0.48300576],
       [0.50364772, 0.49635228],
       [0.51953823, 0.48046177],
       [0.50404011, 0.49595989],
       [0.49833478, 0.50166522],
       [0.5372498 , 0.4627502 ],
       [0.51029376, 0.48970624],
       [0.50165227, 0.49834773],
       [0.49767985, 0.50232095],
       [0.53805393, 0.46194607],
       [0.53085802, 0.46914198],
       [0.49680953, 0.50319047],
       [0.48384828, 0.51615172],
       [0.541238 , 0.458762 ],
       [0.51133426, 0.48866574],
       [0.52415747, 0.47584253],
       [0.47923057, 0.52076943],
       [0.49918054, 0.50081946],
       [0.52934951, 0.47065049],
       [0.49801967, 0.50198033]])

```



File Edit View Insert Cell Kernel Widgets Help

Not Trusted Python 3 (ipykernel)

```
In [15]: yhat_prob.shape
```

```
(40, 2)
```

### Confusion matrix

Another way of looking at accuracy of classifier is to look at confusion matrix.

```
In [16]: from sklearn.metrics import classification_report, confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes, normalize=False, title='Confusion matrix', cmap=plt.cm.Blue)
'''
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting 'normalize=True'.
'''
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

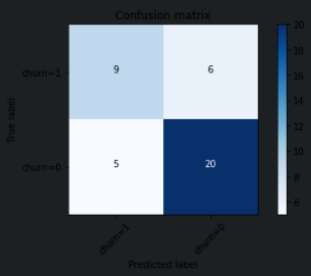
    fmt = '2f' if normalize else 'd'
```

```
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
print(confusion_matrix(y_test, yhat, labels=[1,0]))

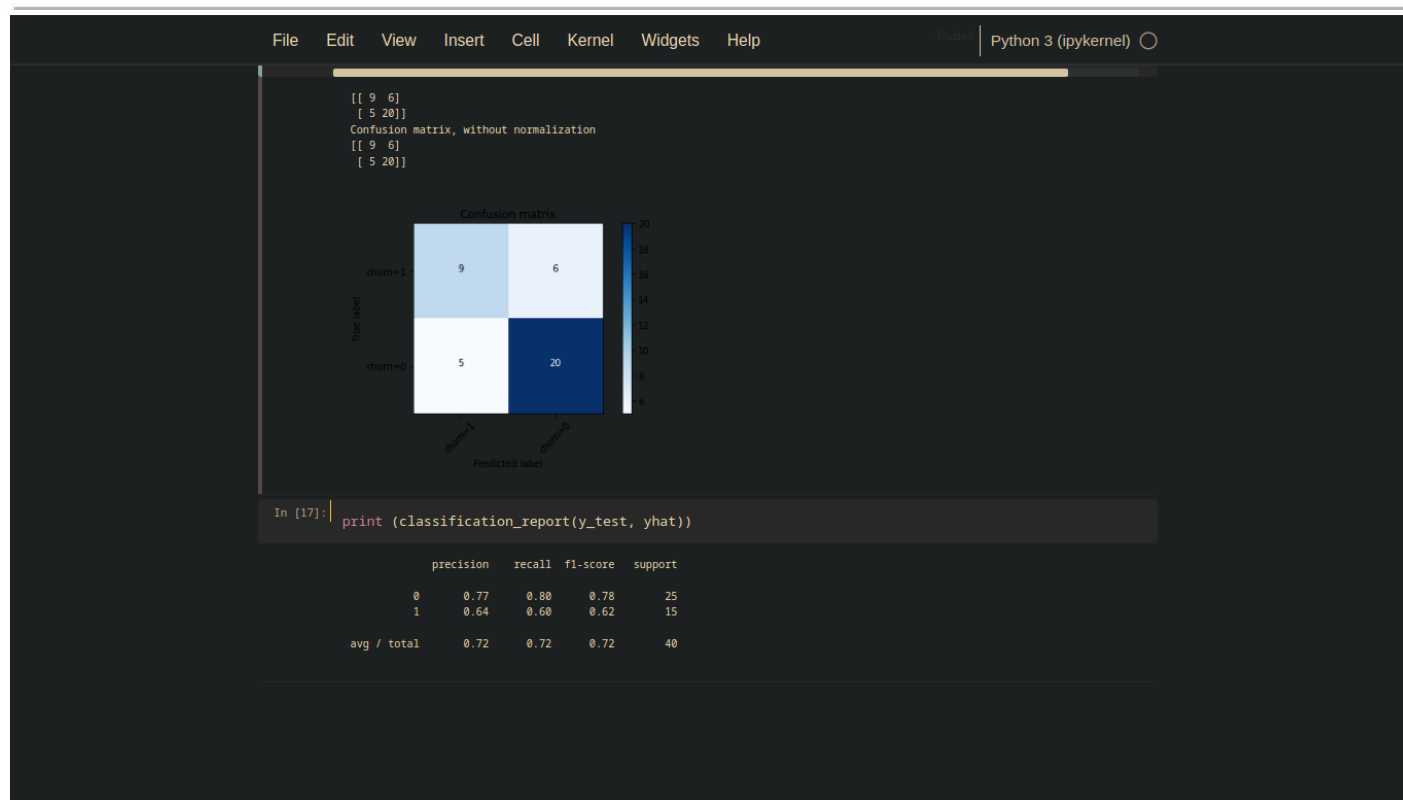
cnf_matrix = confusion_matrix(y_test, yhat, labels=[1,0])
np.set_printoptions(precision=2)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['churn=1','churn=0'],normalize= False, title='Confusio
```

```
[[ 9  6]
 [ 5 20]]
Confusion matrix, without normalization
[[ 9  6]
 [ 5 20]]
```




	churn=1	churn=0
churn=1	9	6
churn=0	5	20





## 5. Observations :



## Learning Outcomes :

- Logistic Regression.
- Tuning the model with the appropriate selection of solver.
- I have learnt about EDA
- I have learnt about Data pre-processing and selection
- I have learnt about Training and Testing
- I have learnt about missing value analysis
- I have learnt Confussion Matrix.

S. No.	Parameters	Marks Obtained	Maximum Marks
1.			
2.			
3.			